Lab 1: Watch one refund go round the loop

Topic 1: Agentic Architecture & Orchestration · Task 1.1: Design and implement agentic loops for autonomous task execution

Runs on Simulator 1: Customer Support Resolution Agent · Run in class

Source: Built from exam guide task statement 1.1 (no practice question covers this angle)

The question

A customer writes in: the ceramic lamp in ORD-5501 arrived cracked, and they want a refund. Your agent has its tools and a loop that keeps calling Claude. When should that loop stop?

A. Stop as soon as the model writes something that sounds finished.

B. Continue on tool_use, stop on end_turn, append every result.

Simulate it live

```
# In the browser: open Simulator 1, then Lab 1
cd labs
uv run serve.py

# In Claude Code, from the labs folder
/lab 1

# In a terminal, from simulators/simulator-1-customer-support/
uv run support_agent.py --case CS-01
```

What you should see

Real stop_reason on every turn: tool_use after get_customer, lookup_order, recommend_refund and process_refund, each result appended, then end_turn stops the loop with a $120 refund.

Open `simulators/simulator-1-customer-support/evidence/` and find it yourself.

The answer

B. Continue on tool_use, stop on end_turn, append every result.

Now change one thing

Copy a file in cases/ to a new CS id, change the customer's message and expected disposition, rerun with --case, then score it with evaluate.py.