Lab 48: The author reviewing its own draft misses its own errors

Topic 4: Prompt Engineering & Structured Output · Task 4.6: Design multi-instance and multi-pass review architectures

Runs on Simulator 3: Multi-Agent Research System · Run in class

Source: SkillCertPro practice exam Part 1, Q6 (unofficial third-party question)

Commit to an answer first

During evaluation, the final reports are well written but sometimes include unsupported claims and miss contradictions between sections. The current design asks the report generator to review its own draft before publishing. What architecture change would most effectively improve review reliability?

A. Delegate verification to listed subagents without passing source excerpts, relying on shared session context to expose prior findings.

B. Have the synthesis agent merge all subagent outputs once, then instruct it to mark unsupported claims before publishing.

C. Ask the report generator to perform an extended self-review immediately after drafting, using its existing conversation history.

D. Run independent reviewer agents on section-specific claim–source bundles, then run a separate integration pass for conflicts and coverage gaps.

Simulate the solution

```
# In the browser: open Simulator 3, then Lab 48
cd labs
uv run serve.py

# In Claude Code, from the labs folder
/lab 48

# In a terminal, from simulators/simulator-3-multi-agent-research/
uv run research_system.py --case UC-06
```

What you should see

Two or more reviewers launch in one response, and a separate integration pass follows.

At least two reviewer subagents launch in the same coordinator response after synthesis, followed by a separate integration pass; the report keeps the review verdicts.

Open `simulators/simulator-3-multi-agent-research/evidence/` and find it yourself.

The answer

D. Independent reviewers on claim-source bundles, then an integration pass.

Separate generation from review. Self-review catches obvious problems but rarely challenges the assumptions built into its own draft. Independent reviewers bring a fresh context, and multi-pass review keeps each pass focused: first local claim-to-source alignment, then cross-section consistency and coverage. Because subagents have isolated context, the coordinator must pass in the claims, excerpts, source URLs and document names. Parallel review suits independent checks, while contradictions only appear in a later integration pass that compares sections.

Why the others fail

A. There is no shared session context to rely on. Reviewers without the excerpts have nothing to verify against.

B. A single merged pass by the agent that produced the synthesis adds no independence and dilutes attention across everything at once.

C. An extended self-review keeps the same reasoning context, along with its blind spots.

Check yourself

- reviewer agents spawned >=2 in one response
- an integration reviewer runs after the section reviewers
- report still includes the coverage gap and contested findings

Now change one thing

In fixtures/sources.json replace the visual-arts TIMEOUT with [] and rerun: a valid empty result must not read as a failure.

Third-party practice material from SkillCertPro, reconstructed from a video transcript in the Tertiary Markdown Vault. Not a real exam item. Wording may differ from the original.