Lab 59: Acceptance looks great. Two fields don't.

Topic 5: Context Management & Reliability · Task 5.5: Design human review workflows and confidence calibration

Runs on Simulator 2: Code Generation with Claude Code · Self-practice

Source: SkillCertPro practice exam Part 1, Q33 (unofficial third-party question)

The question

This team has a custom documentation skill [unclear in transcript] (a slash-invoked command) that updates README sections, API references, and migration notes after refactors. The command's overall acceptance rate is high, but review data shows that incorrect outputs are concentrated in a few fields, especially function signatures and migration warnings. What should the team do before reducing human review?

A. Approve documentation changes when the command's file-level confidence exceeds a fixed threshold.

B. Require field-level confidence outputs and calibrate review thresholds using labeled validation examples.

C. Switch the command to plan mode for all documentation updates before applying changes.

D. Use aggregate acceptance rate as the release gate and sample merged outputs weekly.

Simulate the solution

```
# In the browser: open Simulator 2, then Lab 59
cd labs
uv run serve.py

# In Claude Code, from the labs folder
/lab 59

# In a terminal, from simulators/simulator-2-code-generation/
uv run casekit.py --case UC-33
```

What you should see

One confidence per field. Under 0.8 goes to review, as do signatures under 0.95.

The documentation update returns one confidence per field. The runner routes fields under 0.8 (and every signature or migration warning under 0.95) to human review instead of trusting a file-level score.

Open `simulators/simulator-2-code-generation/evidence/` and find it yourself.

The answer

B. Field-level confidence, calibrated on labelled examples.

For reliability-sensitive generated docs, the field is often the right unit of review. A README can be mostly right and still contain a wrong function signature or an unsafe migration note, so an aggregate metric gives a false sense of readiness. Have Claude output field-level confidence and calibrate it against a labelled validation set. Calibration turns a self-report into a routing signal: low-confidence signatures and migration warnings go to humans, while well-calibrated low-risk fields proceed. This is a review-workflow design problem, not a prompting-mode problem.

Why the others fail

A. File-level confidence hides field-level errors, and a fixed uncalibrated threshold is arbitrary.

C. Plan mode does not provide segmented validation or calibrated thresholds.

D. The aggregate rate is exactly the metric that masks the concentrated errors.

Check yourself

- every field carries a confidence in [0,1]
- the signature field is present for calibration against labelled examples
- learner explains why a file-level threshold or aggregate acceptance rate hides field errors

Now change one thing

Set LAB_MODEL=claude-sonnet-5-5, rerun, and compare the result and the cost line in evidence/.

Third-party practice material from SkillCertPro, reconstructed from a video transcript in the Tertiary Markdown Vault. Not a real exam item. Wording may differ from the original.