Lab 60: The aggregate masks a segment: break it and watch

Topic 5: Context Management & Reliability · Task 5.5: Design human review workflows and confidence calibration

Runs on Simulator 6: Structured Data Extraction · Run in class

Source: Built from exam guide task statement 5.5 (no practice question covers this angle)

Commit to an answer first

Overall accuracy looks excellent, and someone proposes cutting human review. What do you check first?

A. Accuracy by document type and by field.

B. One overall accuracy number.

Simulate the solution

```
# In the browser: open Simulator 6, then Lab 60
cd labs
uv run serve.py

# In Claude Code, from the labs folder
/lab 60

# In a terminal, from simulators/simulator-6-structured-extraction/
uv run extract.py
uv run python -m json.tool evidence/accuracy.json
```

What you should see

Change one invoice total: only by_document_type.invoice falls.

Open `simulators/simulator-6-structured-extraction/evidence/` and find it yourself.

The answer

A. Accuracy by document type and by field.

Now change one thing

Change DOC-04's printed total in fixtures/invoices.txt, rerun, and watch by_document_type.invoice fall while the others hold.