Lab 61: A vendor benchmark isn't a peer-reviewed study

Topic 5: Context Management & Reliability · Task 5.6: Preserve information provenance and handle uncertainty in multi-source synthesis

Runs on Simulator 3: Multi-Agent Research System · Run in class

Source: SkillCertPro practice exam Part 1, Q24 (unofficial third-party question)

Commit to an answer first

During evaluation, the final report cites sources but sometimes treats a peer-reviewed study, a vendor benchmark, and an internal survey as equivalent evidence for the same claim. The search and document-analysis subagents currently return concise finding summaries plus a general list of sources consulted. What change best supports accurate downstream synthesis?

A. Add a report-generation pass that expands the final narrative with a separate bibliography and citation appendix.

B. Require each subagent to return structured claim records with source location, publication date, evidence excerpt, and methodology notes.

C. Ask the coordinator to retain full subagent transcripts and have the synthesizer infer provenance from prior work.

D. Give each subagent broader tool access so it can inspect all available sources before summarizing its findings.

Simulate the solution

```
# In the browser: open Simulator 3, then Lab 61
cd labs
uv run serve.py

# In Claude Code, from the labs folder
/lab 61

# In a terminal, from simulators/simulator-3-multi-agent-research/
uv run research_system.py --case UC-24
```

What you should see

Each finding carries document, page, date and method note. 41% and 55% stay side by side.

Structured claim records, not prose summaries, come back from document-analysis.

Open `simulators/simulator-3-multi-agent-research/evidence/` and find it yourself.

The answer

B. Claim records with location, date, excerpt and method.

Provenance is an interface contract. Downstream agents can rely only on what is explicitly returned in the handoff. Non-fork subagents start with a fresh context and don't share the parent's history. Structured claim records (claim, evidence excerpt, URL or document location, publication or collection date, methodology notes) stop the synthesiser from treating incomparable evidence as interchangeable. This is not just a citation-formatting preference, and it follows the guidance to return high-signal fields.

Why the others fail

A. Formatting at the end cannot recover distinctions already lost upstream.

C. Full transcripts bloat the context, and inferring provenance is unreliable.

D. Inspecting more sources does not help when the handoff still drops the metadata.

Check yourself

- document-analysis result parses as JSON with a findings list
- every finding has document, page, published and methodology_note
- the two SME adoption findings are both present with conflict_with set

Now change one thing

In fixtures/sources.json replace the visual-arts TIMEOUT with [] and rerun: a valid empty result must not read as a failure.

Third-party practice material from SkillCertPro, reconstructed from a video transcript in the Tertiary Markdown Vault. Not a real exam item. Wording may differ from the original.